

Apache Hadoop & MapReduce

Assignment 1


Name: Arpita Kundu, Roll no: M25DE1004

Github link: <https://github.com/arpita12-bot/Machine-learning-with-Big-Data/tree/main>

Apache Hadoop & MapReduce


Q-1:

Run and show the working of the WordCount example as shown on the Apache Hadoop website as well as discussed in class earlier. (No need to send code for this. Just output to show that it is working is enough).



```
arpkundu@DESKTOP-J1UA35:~$ nano wordcount_input.txt
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -mkdir -p /input
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -put wordcount_input.txt /input
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -ls /input
Found 1 items
-rw-r--r-- 1 arpkundu supergroup          63 2026-02-08 22:23 /input/wordcount_input.txt
arpkundu@DESKTOP-J1UA354:~$
```

Wordcount program:



```
GNU nano 6.2                                     WordCount.java
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
                       ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class Reducer
        extends Reducer<Text, IntWritable, Text, IntWritable>{

        public void reduce(Text key, Iterable<IntWritable> values,
                          Context context) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }
}
```

^G Help ^O Write Out ^W Where Is ^K Cut ^T Execute ^C Location M-U Undo M-A Set Mark
 ^X Exit ^R Read File ^\ Replace ^U Paste ^J Justify ^/ Go To Line M-E Redo M-G Copy

Apache Hadoop & MapReduce

Prepare input in HDFS

```
arpkundu@DESKTOP-J1UA354: ~$ nano wc.txt
GNU nano 6.2
hello hadoop
hello mapreduce
hello big data
^G Help      ^O Write Out    ^W Where Is      ^K Cut
^X Exit      ^R Read File    ^N Replace      ^U Paste
^T Execute   ^J Justify     ^C Location     ^/ Go To Line
^M-U Undo    ^M-E Redo      ^M-A Set Mark
^M-G Copy
```

Output:

```
arpkundu@DESKTOP-J1UA354: ~$ hdfs dfs -cat /wordcount/output/part-r-00000
big      1
data     1
hadoop   1
hello    3
mapreduce      1
```

Q-2:

Input file we use with following lyrics named as wc1.txt.

```
We're up all night till the sun
We're up all night to get some
We're up all night for good fun
We're up all night to get lucky
```

The input pairs for the Map phase will be the following:

```
(0, "We're up all night to the sun")
(31, "We're up all night to get some")
(63, "We're up all night for good fun")
(95, "We're up all night to get lucky")
```

The key is the byte offset starting from the beginning of the file. While we won't need this value in Word Count, it is always passed to the Mapper by the Hadoop framework.

Apache Hadoop & MapReduce

The byte offset is a number that can be large if there are many lines in the file.

```
arpkundu@DESKTOP-J1UA35: ~$ nano wc1.txt
GNU nano 6.2
We're up all night till the sun
We're up all night to get some
We're up all night for good fun
We're up all night to get lucky

^C Help      ^O Write Out    ^W Where Is      ^K Cut
^X Exit      ^R Read File    ^\ Replace      ^U Paste
^T Execute   ^J Justify     ^G Location    ^/ Go To Line
M-U Undo     M-E Redo      M-A Set Mark
M-G Copy
```

Put the input file into HDFS using following commands:

```
hdfs dfs -mkdir -p /lyrics/input
```

```
hdfs dfs -put wc1.txt /lyrics/input
```

```
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -mkdir -p /lyrics/input
hdfs dfs -put wc1.txt /lyrics/input
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -ls /lyrics/input
Found 1 items
-rw-r--r--  1 arpkundu supergroup          129 2026-02-08 23:18 /
lyrics/input/wc1.txt
arpkundu@DESKTOP-J1UA354:~$ hadoop jar $HADOOP_HOME/share/hadoop
```

Now run wordcount:

```
hadoop jar $HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
wordcount /lyrics/input /lyrics/output
```

Input to Map phase:

```
(LongWritable, Text)
(0, "We're up all night till the sun")
(31, "We're up all night to get some")
(63, "We're up all night for good fun")
(95, "We're up all night to get lucky")
```

Result :

```
hdfs dfs -cat /lyrics/output/part-r-00000
```

Apache Hadoop & MapReduce

```
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -cat /lyrics/output/part-r-00000
We're    4
all      4
for      1
fun      1
get      2
good     1
lucky    1
night   4
some    1
sun     1
the     1
till    1
to      2
up      4
arpkundu@DESKTOP-J1UA354:~$
```

Explanation : The WordCount MapReduce program was executed on Hadoop using the given lyrics as input. The output shows correct word frequencies, confirming successful execution of the Map and Reduce phases.

1. What will the output pairs look like(Mapper)?

In **WordCount**, the **Mapper**:

- Reads **one line at a time**
- Splits the line into words
- Emits **(word, 1)** for each word

Output pairs from the Map phase:

Each word generates one key-value pair:

```
(We're, 1)
(up, 1)
(all, 1)
(night, 1)
(till, 1)
(the, 1)
(sun, 1)
```

These are the **output pairs of the Map phase, not reduced yet.**

2. What are the **types of keys and values** in the Map phase?

Input to the Mapper :

Component	Type (Hadoop)	Meaning
Key	LongWritable	Byte offset of the line in the file
Value	Text	One line of input text

Apache Hadoop & MapReduce

Example:

(LongWritable, Text)
 (0, "We're up all night till the sun")

Output from the Mapper:

Component	Type (Hadoop)	Meaning
Key	Text	A word
Value	IntWritable	Count = 1

Example:

(Text, IntWritable)
 ("night", 1)

Q-3:

For the reduce phase, some of the output pairs will be the following:

("up", 4)
 ("to", 3)
 ("get", 2)
 ("lucky", 1)

1. What will the input pairs look like?

Before the Reduce phase, Hadoop performs the **Shuffle and Sort** step.
 This step **groups all values with the same key** together.

So, the reducer receives input in the form:

(word, list_of_values)

Reduce Phase Input Pairs:

("up", [1, 1, 1, 1])
 ("to", [1, 1, 1])
 ("get", [1, 1])
 ("lucky", [1])

Each 1 represents one occurrence of the word emitted by the Mapper.

2. What will be the types of keys and values of the input and output pairs in the Reduce phase?

Input to the Reducer:

Component	Hadoop Data Type	Description
Key	Text	A word
Value	Iterable<IntWritable>	List of counts

Example:

(Text, Iterable<IntWritable>)
 ("up", [1, 1, 1, 1])

Output from the Reducer:

Component	Hadoop Data Type	Description
Key	Text	A word
Value	IntWritable	Total count

Example:

(Text, IntWritable)
 ("up", 4)

In the Reduce phase, input pairs consist of a word and an iterable list of counts, and output pairs consist of the word and its total frequency. Reducer receives grouped values and produces final word counts.

Q-4:

See the **WordCount.java** code from the MapReduce Tutorial, and copy its contents to the corresponding file in your project. Find the definitions of the Map class and the map() function:


```
public static class Map extends Mapper {
@Override
public void map(/*?*/ key, /*?*/ value, Context context)
throws IOException, InterruptedException {
...
}
```

Use the data types you found in Question 1 to replace the /*?*/ placeholders.

Similarly, find the definitions of the Reduce class and the reduce() function and replace the /*?*/ placeholders with the data types found in Question 2.

You will also have to find which arguments to pass to job.setOutputKeyClass() and job.setOutputValueClass() .

Apache Hadoop & MapReduce




```

GNU nano 6.2                                     WordCount.java
import java.io.IOException;
import java.util.StringTokenizer;
|
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {
    public static class Map
        extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context)
            throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
[ Read 63 lines ]
^G Help      ^O Write Out   ^W Where Is   ^K Cut       ^T Execute   ^C Location   M-U Undo
^X Exit      ^R Read File   ^\ Replace    ^U Paste     ^J Justify   ^/ Go To Line M-E Redo
                                         M-A Set Mark
                                         M-G Copy

```



```


GNU nano 6.2                                     WordCount.java
import java.io.IOException;
import java.util.StringTokenizer;
|
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {
    public static class Map
        extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context)
            throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
[ Read 63 lines ]
^G Help      ^O Write Out   ^W Where Is   ^K Cut       ^T Execute   ^C Location   M-U Undo
^X Exit      ^R Read File   ^\ Replace    ^U Paste     ^J Justify   ^/ Go To Line M-E Redo
                                         M-A Set Mark
                                         M-G Copy

```

Apache Hadoop & MapReduce



```

GNU nano 6.2                                     WordCount.java
throws IOException, InterruptedException {
    StringTokenizer itr = new StringTokenizer(value.toString());
    while (itr.hasMoreTokens()) {
        word.set(itr.nextToken());
        context.write(word, one);
    }
}


public static class Reduce
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
    }
}

^G Help      ^O Write Out   ^W Where Is   ^K Cut       ^T Execute   ^C Location   M-U Undo   M-A Set Mark
^X Exit      ^R Read File   ^\ Replace    ^U Paste     ^J Justify   ^/ Go To Line M-E Redo   M-6 Copy

```



```

GNU nano 6.2                                     WordCount.java
throws IOException, InterruptedException {
    StringTokenizer itr = new StringTokenizer(value.toString());
    while (itr.hasMoreTokens()) {
        word.set(itr.nextToken());
        context.write(word, one);
    }
}

public static class Reduce
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
    }
}

^G Help      ^O Write Out   ^W Where Is   ^K Cut       ^T Execute   ^C Location   M-U Undo   M-A Set Mark
^X Exit      ^R Read File   ^\ Replace    ^U Paste     ^J Justify   ^/ Go To Line M-E Redo   M-6 Copy

```

Apache Hadoop & MapReduce

```
arpkundu@DESKTOP-J1UA35: ~ + v
GNU nano 6.2                                     WordCount.java

    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

Q-5:

Write the map() function. We want to make sure to disregard punctuation: to this end, you can use String.replaceAll(). In order to split lines into words, you can use a StringTokenizer .

```
public class WordCount {

    public static class Map
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        @Override
        protected void map(Object key, Text value, Context context)
            throws IOException, InterruptedException {
            String line = value.toString()
                .toLowerCase()
                .replaceAll("[^a-z0-9\\s]", "");
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

Apache Hadoop & MapReduce

```

public class WordCount {

    public static class Map
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        @Override
        protected void map(Object key, Text value, Context context)
            throws IOException, InterruptedException {

            String line = value.toString()
                .toLowerCase()
                .replaceAll("[^a-zA-Z0-9\\s]", "");

            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }
}

```

Q-6:

Write the reduce() function. When you're done, make sure that compiling the project doesn't produce any errors.

```

public static class Reduce
    extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    protected void reduce(Text key, Iterable<IntWritable> values,
        Context context)
        throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

```

Apache Hadoop & MapReduce

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");

    job.setJarByClass(WordCount.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Run the Program without errors

```

arpkundu@DESKTOP-J1UA354:~$ ls
'WordCount$IntSumReducer.class'      mlbd
'WordCount$TokenizerMapper.class'    mlbd.pub
WordCount.class                      wc.txt
WordCount.java                       wc1.txt
authorized_keys                      wordcount.jar
authorized_keys.pub                  wordcount_input.txt
hadoop-3.3.6.tar.gz

arpkundu@DESKTOP-J1UA354:~$ javac -classpath `hadoop classpath` WordCount.java

arpkundu@DESKTOP-J1UA354:~$ ls
'WordCount$IntSumReducer.class'      hadoop-3.3.6.tar.gz
'WordCount$Map.class'                mlbd
'WordCount$Reduce.class'             mlbd.pub
'WordCount$TokenizerMapper.class'   wc.txt
WordCount.class                     wc1.txt
WordCount.java                      wordcount.jar
authorized_keys                     wordcount_input.txt
authorized_keys.pub

```

Apache Hadoop & MapReduce

```
arpkundu@DESKTOP-J1UA354:~$ hdfs dfs -cat /wc/output/part-r-00000
all      4
for      1
fun      1
get      2
good     1
lucky    1
night    4
some     1
sun      1
the      1
till     1
to       2
up       4
were     4
arpkundu@DESKTOP-J1UA354:~$
```

Q-7:

We will now run WordCount on the 200.txt file.

```
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ ls
10.txt 13.txt 160.txt 206.txt 234.txt 263.txt 291.txt 318.txt 348.txt 375.txt 407.txt 44.txt 469.txt 496.txt 77.txt
101.txt 130.txt 161.txt 207.txt 235.txt 264.txt 292.txt 32.txt 349.txt 376.txt 408.txt 440.txt 47.txt 498.txt 78.txt
102.txt 131.txt 162.txt 208.txt 236.txt 265.txt 293.txt 321.txt 35.txt 377.txt 409.txt 442.txt 470.txt 499.txt 79.txt
103.txt 132.txt 163.txt 209.txt 237.txt 266.txt 295.txt 322.txt 350.txt 378.txt 41.txt 443.txt 471.txt 5.txt 8.txt
104.txt 133.txt 164.txt 21.txt 238.txt 267.txt 296.txt 323.txt 351.txt 379.txt 410.txt 444.txt 472.txt 50.txt 80.txt
105.txt 134.txt 165.txt 210.txt 239.txt 268.txt 297.txt 324.txt 352.txt 38.txt 411.txt 445.txt 473.txt 51.txt 81.txt
106.txt 136.txt 166.txt 211.txt 24.txt 27.txt 298.txt 326.txt 353.txt 380.txt 412.txt 446.txt 474.txt 52.txt 82.txt
107.txt 137.txt 167.txt 212.txt 240.txt 270.txt 299.txt 327.txt 354.txt 381.txt 414.txt 447.txt 475.txt 53.txt 83.txt
108.txt 138.txt 168.txt 213.txt 241.txt 271.txt 3.txt 328.txt 355.txt 382.txt 416.txt 448.txt 476.txt 54.txt 85.txt
109.txt 139.txt 169.txt 214.txt 242.txt 272.txt 30.txt 33.txt 356.txt 383.txt 417.txt 449.txt 477.txt 55.txt 87.txt
11.txt   14.txt 17.txt 215.txt 243.txt 273.txt 300.txt 330.txt 357.txt 384.txt 418.txt 45.txt 478.txt 56.txt 88.txt
110.txt 140.txt 170.txt 216.txt 244.txt 274.txt 301.txt 331.txt 358.txt 385.txt 419.txt 450.txt 479.txt 57.txt 9.txt
111.txt 143.txt 171.txt 217.txt 246.txt 275.txt 302.txt 332.txt 359.txt 389.txt 42.txt 451.txt 48.txt 58.txt 90.txt
112.txt 144.txt 172.txt 219.txt 249.txt 276.txt 303.txt 333.txt 360.txt 39.txt 420.txt 452.txt 480.txt 6.txt 92.txt
113.txt 145.txt 173.txt 22.txt 25.txt 277.txt 304.txt 334.txt 361.txt 390.txt 421.txt 453.txt 481.txt 60.txt 94.txt
114.txt 146.txt 174.txt 220.txt 250.txt 278.txt 305.txt 335.txt 362.txt 391.txt 422.txt 454.txt 482.txt 61.txt 96.txt
115.txt 148.txt 175.txt 221.txt 251.txt 279.txt 306.txt 336.txt 363.txt 392.txt 423.txt 455.txt 483.txt 62.txt 97.txt
117.txt 149.txt 176.txt 222.txt 252.txt 28.txt 307.txt 337.txt 364.txt 396.txt 424.txt 456.txt 484.txt 63.txt 98.txt
118.txt 15.txt 178.txt 223.txt 253.txt 280.txt 308.txt 338.txt 365.txt 397.txt 427.txt 457.txt 485.txt 64.txt 99.txt
12.txt   150.txt 179.txt 224.txt 254.txt 281.txt 309.txt 339.txt 366.txt 398.txt 428.txt 458.txt 486.txt 65.txt
121.txt 151.txt 18.txt 226.txt 255.txt 283.txt 31.txt 34.txt 367.txt 399.txt 429.txt 459.txt 487.txt 66.txt
122.txt 152.txt 180.txt 227.txt 256.txt 284.txt 310.txt 340.txt 368.txt 4.txt 430.txt 46.txt 488.txt 67.txt
123.txt 153.txt 19.txt 228.txt 257.txt 285.txt 311.txt 341.txt 369.txt 400.txt 431.txt 460.txt 489.txt 68.txt
124.txt 154.txt 2.txt 229.txt 258.txt 286.txt 312.txt 342.txt 37.txt 401.txt 432.txt 461.txt 49.txt 69.txt
125.txt 156.txt 20.txt 23.txt 259.txt 287.txt 313.txt 343.txt 370.txt 402.txt 433.txt 462.txt 490.txt 7.txt
126.txt 157.txt 200.txt 230.txt 26.txt 288.txt 314.txt 344.txt 371.txt 403.txt 434.txt 463.txt 491.txt 71.txt
127.txt 158.txt 201.txt 231.txt 260.txt 289.txt 315.txt 345.txt 372.txt 404.txt 436.txt 465.txt 493.txt 72.txt
128.txt 159.txt 202.txt 232.txt 261.txt 29.txt 316.txt 346.txt 373.txt 405.txt 437.txt 466.txt 494.txt 73.txt
129.txt 16.txt 204.txt 233.txt 262.txt 290.txt 317.txt 347.txt 374.txt 406.txt 439.txt 467.txt 495.txt 75.txt
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ pwd
/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ ls | grep "200.txt"
200.txt
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$
```

Copy 200.txt into HDFS

```
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ hdfs dfs -copyFromLocal 200.txt /user/arpkundu
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ hdfs dfs -ls /user/arpkundu
Found 1 items
-rw-r--r-- 1 arpkundu supergroup 8312639 2026-02-10 19:22 /user/arpkundu/200.txt
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$
```

```

arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ hdfs dfs -ls output
-rw-r--r-- 1 arpkundu supergroup 0 2026-02-10 19:25 output/_SUCCESS
-rw-r--r-- 1 arpkundu supergroup 863649 2026-02-19:25 output/part-r-00000
arpkundu@DESKTOP-J1UA354:/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB$ Found 2 items
zurbriggen 1 zurbriggen 1
zurich 15 zurich 15
zurichgrebel 1 zurichgrebel 1
zurita 1 zurita 1
zurmat 1 zurmat 1
zusammenhang 1 zusammenhang 1
zutphen 1 1 zutphen 1
zweckmasigkeit 1 zweckmasigkeit 1
zwei 4 1 zwei 4
zweite 1 zweite 1
zweyer 1 zwolkau 2
zwickau 2 zwinglias 1
zwinglians 1 zygoma 1
zwinglius 1 zyryanovsk 1
zwolf 1 zygote 3

```

Q-8:**Why don't we have a replication factor for directories in HDFS?**

In HDFS, **replication applies only to file blocks**, not directories.

1. Directories Do Not Store Data Blocks:

- Files in HDFS are split into **blocks**.
- Each block is replicated (e.g., 3 copies).
- Directories **do not contain data blocks**.
- They only store **metadata** (like file names and structure).

Since replication is about copying **data blocks across DataNodes**, directories don't need replication.

2. Directory Information is Stored in the NameNode

The **NameNode** manages:

- Directory structure
- File names
- Block locations

This metadata is stored centrally (with FSImage and EditLogs).

It is not distributed across DataNodes like file blocks.

So directories are already managed through metadata — not through block replication.

3. Replicating Directories Would Be Redundant

If HDFS replicated directories:

- It would duplicate only metadata references
- But the actual data replication already happens at block level
- So directory-level replication would serve no purpose

Directories in HDFS do not have a replication factor because replication applies only to file blocks. Directories contain only metadata, which is managed by the NameNode. Since no data blocks are stored in directories, replication is unnecessary.

Apache Hadoop & MapReduce

Q-9:


Edit WordCount.java to make it measure and display the total execution time of the job.

Experiment with the mapreduce.input.fileinputformat.split.maxsize parameter.

You can change its value using:


```
job.getConfiguration().setLong("mapreduce.input.fileinputformat.split.maxsize", );
```

PART 1 : Measure and display total execution time



```
arpkundu@DESKTOP-J1UA354:~$ nano WordCount.java
GNU nano 6.2                               WordCount.java


public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    // ----- START TIME -----
    long startTime = System.currentTimeMillis();
    boolean success = job.waitForCompletion(true);
    // ----- END TIME -----
    long endTime = System.currentTimeMillis();
    long totalTime = endTime - startTime;
    System.out.println("=====Total Job Execution Time (ms): " + totalTime);
}
^G Help      ^O Write Out   ^W Where Is   ^K Cut          ^T Execute   ^C Location   M-U Undo   M-A Set Mark
^X Exit      ^R Read File   ^\ Replace    ^U Paste       ^J Justify   ^/ Go To Line M-E Redo   M-G Copy
```



```
arpkundu@DESKTOP-J1UA354:~$ ls
200.txt          'WordCount$TokenizerMapper.class'  authorized_keys.pub  wc.txt
'WordCount$IntSumReducer.class'  WordCount.class        hadoop-3.3.6.tar.gz  wcl.txt
'WordCount$Map.class'           WordCount.java         mldb                   wordcount.jar
'WordCount$Reduce.class'        authorized_keys       mldb.pub                wordcount_input.txt
arpkundu@DESKTOP-J1UA354:~$ javac -classpath `hadoop classpath` WordCount.java
arpkundu@DESKTOP-J1UA354:~$ nano WordCount.java
arpkundu@DESKTOP-J1UA354:~$ ls
200.txt          'WordCount$TokenizerMapper.class'  authorized_keys.pub  wc.txt
'WordCount$IntSumReducer.class'  WordCount.class        hadoop-3.3.6.tar.gz  wcl.txt
'WordCount$Map.class'           WordCount.java         mldb                   wordcount.jar
'WordCount$Reduce.class'        authorized_keys       mldb.pub                wordcount_input.txt
arpkundu@DESKTOP-J1UA354:~$ jar cf wordcount.jar WordCount*.class
arpkundu@DESKTOP-J1UA354:~$ ls | grep wordcount.jar
wordcount.jar
arpkundu@DESKTOP-J1UA354:~$
```

PART 2: Experiment with split.maxsize

Apache Hadoop & MapReduce



```

arpkundu@DESKTOP-JIU354: ~      +  ×
GNU nano 6.2                         WordCount.java

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");

        job.setJarByClass(WordCount.class);
        job.setMapperClass(Map.class);
        job.setCombinerClass(Reduce.class);
        job.setReducerClass(Reduce.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        job.getConfiguration().setLong("mapreduce.input.fileinputformat.split.maxsize",
67108864);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        // ----- START TIME -----
        long startTime = System.currentTimeMillis();
        boolean success = job.waitForCompletion(true);

        // ----- END TIME -----
        long endTime = System.currentTimeMillis();

    ^G Help      ^O Write Out   ^W Where Is   ^K Cut       ^T Execute   ^C Location   M-U Undo
    ^X Exit      ^R Read File   ^\ Replace    ^U Paste     ^J Justify   ^/ Go To Line M-E Redo
                                         ^A Set Mark   M-G Copy


```

Compile : javac -classpath `hadoop classpath` WordCount.java

Create jar:

jar cf wordcount.jar WordCount*.class

Run the job:



```

arpkundu@DESKTOP-JIU354: ~ $ hadoop jar wordcount.jar WordCount /user/arpkundu/200.txt output
2026-02-10 19:53:32,121 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2026-02-10 19:53:32,211 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2026-02-10 19:53:32,211 INFO impl.MetricsSystemImpl: JobTracker metrics system started
Exception in thread "main" org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory hdfs://localhost:9000/user/arpkundu/output already exists
    at org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileOutputFormat.java:164)
    at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:277)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:143)
    at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1678)
    at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1675)
    at java.base/java.security.AccessController.doPrivileged(Native Method)
    at java.base/javax.security.auth.Subject.doAs(Subject.java:423)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1899)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1675)
    at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1696)
    at WordCount.main(WordCount.java:77)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.base/java.lang.reflect.Method.invoke(Method.java:566)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:328)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:241)
arpkundu@DESKTOP-JIU354: ~ $

```

How does changing mapreduce.input.fileinputformat.split.maxsize impact performance? Why?

mapreduce.input.fileinputformat.split.maxsize

- controls the **maximum size of input splits**, which determines how many **Mapper tasks** are created.

What Happens When You Change It?

Smaller Split Size (e.g., 32MB)

- Creates more input splits
- More Map tasks

If too small → performance may decrease due to overhead.

Larger Split Size (e.g., 256MB or 512MB)

- Creates fewer input splits
- Fewer Map tasks

If too large → performance may decrease due to underutilization.

Why This Happens?

Hadoop processes data **per split**, and:

- 1 split = 1 Mapper
- More splits = more parallel processing
- But too many small tasks = scheduling + context switching overhead
- Too few large tasks = not enough parallelism

So performance depends on **balancing parallelism vs overhead**.

Ex-

Suppose:

- File size = 1GB
- Default block size = 128MB

Default:

$1\text{GB} / 128\text{MB} \approx 8$ Mappers

If `split.maxsize = 32MB`:

$1\text{GB} / 32\text{MB} \approx 32$ Mappers
More parallel but more overhead

If `split.maxsize = 512MB`:

$1\text{GB} / 512\text{MB} \approx 2$ Mappers
Less overhead but low parallelism

Usually:

- Keep split size close to **HDFS block size**
- Tune based on:
 - Cluster size
 - CPU cores
 - Dataset size

Changing `mapreduce.input.fileinputformat.split.maxsize` affects the number of input splits and therefore the number of mapper tasks. Smaller split sizes increase parallelism but add scheduling overhead, while larger split sizes reduce overhead but may underutilize cluster resources. Optimal performance is achieved by balancing parallelism and task management overhead, usually by keeping split size close to the HDFS block size.


Apache Spark

Note: For these Spark problems, first load the provided Project Gutenberg dataset into a Spark DataFrame named `books_df` with the schema:

- **file_name (string):** The name of the text file (e.g., "10.txt").
- **text (string):** The raw text content of the book.

Q-10:

Spark installation completed.



A terminal window showing the command-line steps to install Apache Spark. The terminal is running on a Linux system with the user 'arpkundu'. The steps involve moving the Spark distribution from its default location to a directory named 'spark' and then editing the '.bashrc' file to set the environment variable 'SPARK_HOME'.

```
arpkundu@DESKTOP-J1UA35:~$ mv spark-3.5.1-bin-hadoop3 spark
arpkundu@DESKTOP-J1UA35:~$ nano ~/.bashrc
arpkundu@DESKTOP-J1UA35:~$ source ~/.bashrc
```

Apache Hadoop & MapReduce

Book Metadata Extraction and Analysis:

PART 1: Metadata Extraction

```
>>> from pyspark.sql import Row
>>> from pyspark.sql.functions import regexp_extract
>>> rdd = spark.sparkContext.wholeTextFiles(
...     "/mnt/c/Users/ARPITA KUNDU/Downloads/D184MB/D184MB/*.txt"
... )
>>> books_df = rdd.map(lambda x: Row(
...     file_name=x[0].split("/")[-1],
...     text=x[1]
... )).toDF()
```

file_name	title	release_date	language	encoding
10.txt	The King James Bible	March 2, 2011 [EBook #10]	English	ASCII
101.txt	Hacker Crackdown	January, 1994	English	ASCII
102.txt	The Tragedy of Pudd'nhead Wilson	January, 1994	English	ASCII
103.txt	Around the World in 80 Days	May 15, 2008 [EBook #103]	English	ASCII
104.txt	Franklin Delano Roosevelt's First Inaugural Address	May 14, 2008 [EBook #104]	English	ASCII

PART 2: Analysis

Number of books released each year:

```
>>> metadata_df = metadata_df.withColumn(
...     "release_year",
...     regexp_extract(col("release_date"), r"(\d{4})", 1)
... )
>>> |
```

Apache Hadoop & MapReduce

```
>>> books_per_year = metadata_df \
    .groupBy("release_year") \
    .count() \
    .orderBy("release_year")

books_per_year.show()
...     .filter(col("release_year") != "") \
...     .groupBy("release_year") \
...     .count() \
...     .orderBy("release_year")
>>>
```

```
>>> books_per_year.show()
+-----+---+
|release_year|count|
+-----+---+
| 1975 |    1 |
| 1978 |    1 |
| 1979 |    1 |
| 1991 |    7 |
| 1992 |   19 |
| 1993 |   13 |
| 1994 |   17 |
| 1995 |   60 |
| 1996 |   53 |
| 2002 |    1 |
| 2004 |    7 |
| 2005 |    4 |
| 2006 |   42 |
| 2007 |   13 |
| 2008 |  154 |
| 2009 |    1 |
| 2010 |    9 |
| 2011 |    1 |
| 2012 |    2 |
| 2013 |    1 |
+-----+---+
only showing top 20 rows
```

Most common language in the dataset:

```
>>> metadata_df \
...     .filter(col("language") != "") \
...     .groupBy("language") \
...     .count() \
...     .orderBy(col("count").desc()) \
...     .show(1)
+-----+---+
|language|count|
+-----+---+
| English| 404|
+-----+---+
only showing top 1 row
>>
```

Average length of book titles (characters):

```
>>> metadata_df \
...     .filter(col("title") != "") \
...     .select(length(col("title")).alias("title_length")) \
...     .agg({"title_length": "avg"}) \
...     .show()
+-----+
| avg(title_length)|
+-----+
|22.829268292682926|
+-----+
```

PART 3:

Explain the regular expressions you used to extract the title, release date, language, and encoding. Discuss any challenges or limitations in using regular expressions for this task.

Regular expressions used in task:

Title:

(?i)Title:\s*(.+)

(?i) → case-insensitive

Matches the title line and captures the full title text.

Release Date:

(?i)Release Date:\s*(.+)

Captures the release date line

Year extracted later using \d{4}

Language:

(?i)Language:\s*(.+)

Extracts language metadata (ex- English)

Encoding:

(?i)Character set encoding:\s*(.+)

Extracts text encoding (e.g., UTF-8)

Challenges / Limitations**Some major challenges:**

1. **Missing metadata** : Some books do not contain all fields
2. **Inconsistent formatting** : Metadata labels may vary across books
3. **Multiple matches** : Regex captures only the first occurrence
4. **Noise in text** : Metadata may appear outside the header

PART 4:

What are some potential issues with the extracted metadata (e.g., inconsistencies, missing values)? How would you handle these issues in a real-world scenario?

How to handle in real-world scenarios?

1. Filter empty values.
 2. Normalize text before extraction.
 3. Use structured metadata formats (XML / JSON).
 4. Combine regex with NLP-based extraction.
 5. Apply data validation and defaults.
- Metadata such as title, release date, language, and encoding were extracted from the raw text using regular expressions applied to each book loaded as a single record.
 - Apache Spark's regexp_extract function was used for efficient metadata extraction. The number of books released per year was calculated by extracting the year from the release date. The most common language was identified using frequency analysis, and the average title length was computed using string length aggregation.
 - Regular expressions may fail when metadata is missing or inconsistently formatted; in real-world applications, such issues can be addressed using normalization, validation, and structured metadata sources.

Q-11:**TF-IDF and Book Similarity:**

Task: Calculate TF-IDF scores for words in each book and use them to determine book similarity based on cosine similarity.

PART 1: Preprocessing

Clean the text column in books_df : remove Project Gutenberg header/footer, convert to lowercase, remove punctuation, tokenize into words, and remove stop words.

Data Cleaning:

Remove:

- Gutenberg header/footer (simple heuristic)
- Convert to lowercase
- Remove punctuation
- Tokenize
- Remove stopwords

Import required libraries:

```
Using Python version 3.10.12 (main, Jan 26 2026 14:55:28)
Spark context Web UI available at http://10.255.255.254:4040
Spark context available as 'sc' (master = local[*], app id = local-1770803791182).
SparkSession available as 'spark'.
>>> from pyspark.ml.feature import Tokenizer, StopWordsRemover
>>> from pyspark.sql.functions import col, lower, regexp_replace
>>>
```

file_name	title	release_date	language	encoding
10.txt	The King James Bible	[March 2, 2011 [EBook #10]]	English	ASCII
101.txt	Hacker Crackdown	[January, 1994]	English	ASCII
102.txt	The Tragedy of Pudd'nhead Wilson	[January, 1994]	English	ASCII
103.txt	Around the World in 80 Days	[May 15, 2008 [EBook #103]]	English	ASCII
104.txt	Franklin Delano Roosevelt's First Inaugural Address	[May 14, 2008 [EBook #104]]	English	ASCII
105.txt	Persuasion	[June 5, 2008 [EBook #105]]	English	ASCII
106.txt	Jungle Tales of Tarzan	[June 5, 2008 [EBook #106]]	English	ASCII
107.txt	Far from the Madding Crowd	[February, 1994 [eBook #107]]	English	ISO-646-US (US-ASCII)
108.txt	The Return of Sherlock Holmes	[July 8, 2007 [EBook #108]]	English	ASCII
109.txt	Renascence and Other Poems	[June 19, 2008 [EBook #109]]	English	ASCII

Data Cleaning: convert to lowercase and remove punctuation

```
>>> clean_df = books_df.withColumn(
    "clean_text",
    lower(rege...      "clean_text",
...       lower(regexp_replace(col("text"), "[^a-zA-Z\\s]", " ")))
... )
```

```
>>> clean_df.show()
26/02/11 15:42:13 WARN PythonRunner: Detected deadlock while completing task 0.0
+-----+-----+-----+
|file_name|      text|    clean_text|
+-----+-----+-----+
| 10.txt|The Project Guten...|the project guten...
| 101.txt|The Project Guten...|the project guten...
| 102.txt|The Project Guten...|the project guten...
| 103.txt|The Project Guten...|the project guten...
| 104.txt|The Project Guten...|the project guten...
| 105.txt|The Project Guten...|the project guten...
| 106.txt|Project Gutenberg...|project gutenberg...
| 107.txt|The Project Guten...|the project guten...
| 108.txt|Project Gutenberg...|project gutenberg...
| 109.txt|Project Gutenberg...|project gutenberg...
| 11.txt|Project Gutenberg...|project gutenberg...
| 110.txt|The Project Guten...|the project guten...
| 111.txt|The Project Guten...|the project guten...
| 112.txt|The Project Guten...|the project guten...
| 113.txt|The Project Guten...|the project guten...
| 114.txt|The Project Guten...|the project guten...
| 115.txt|The Project Guten...|the project guten...
| 117.txt|The Project Guten...|the project guten...
| 118.txt|The Project Guten...|the project guten...
| 12.txt|The Project Guten...|the project guten...
+-----+-----+-----+
only showing top 20 rows
```

Tokenize :

```
>>> tokenizer = Tokenizer(inputCol="clean_text", outputCol="words")
rm(clean_df)
>>> tokenized_df = tokenizer.transform(clean_df)
>>> |
```

Remove stop words:

```
>>> remover = StopWordsRemover(inputCol="words", outputCol="filtered_words")
ransform(tokenized_df)
>>> filtered_df = remover.transform(tokenized_df)
>>>
```

PART 2: TF-IDF Calculation

1. Calculate the Term Frequency (TF) of each word in each book.

Import library

```
>>> from pyspark.ml.feature import HashingTF, IDF
>>>
```

Term Frequency (TF):

```
>>> hashingTF = HashingTF(
    Col="raw_feature...      inputCol="filtered_words",
    s",
    numFeatu...      outputCol="raw_features",
    res=10000
)

fea...      numFeatures=10000
hingTF.transform... )
(filtered_df)
>>>
>>> featurized_df = hashingTF.transform(filtered_df)
>>>
```

Inverse Document Frequency (IDF):

2. Calculate the Inverse Document Frequency (IDF) for each word across all books.
3. Compute the TF-IDF score for each word in each book (TF * IDF).

PART 3: Book Similarity

1. Represent each book as a vector of its TF-IDF scores.
2. Calculate the cosine similarity between all pairs of book vectors.
3. For a given book (e.g., file_name "10.txt"), identify the top 5 most similar books based on cosine similarity.

PART 4:

Explain TF and IDF. Why is TF-IDF useful for weighting the importance of words in a document?

What is TF?

- Term Frequency (TF) measures how often a word appears in a document.

$$TF = \frac{\text{word count in document}}{\text{total words in document}}$$

What is IDF?

- Inverse Document Frequency measures how rare a word is across documents.

$$IDF = \log \left(\frac{N}{df} \right)$$

- N = total number of documents
- df = number of documents containing the word

Why is TF-IDF useful for weighting?

Apache Hadoop & MapReduce

- Increases weight of important words and Suppresses common words (e.g., “love”, “the”, “and”)
- Reduces weight of common words
- Highlights distinguishing terms
- Improves document comparison and helps identify document themes
- Improves similarity detection

Suppose we have two books:

Book A (Romance) : love love love marriage happiness love

Book B (Science) : quantum physics atom energy quantum experiment

Step 1: Term Frequency (TF)

In Book A:

- **love** appears 4 times → High TF
- **marriage** appears 1 time

In Book B:

- **quantum** appears 2 times
- **physics** appears 1 time

So TF measures frequency **within the document only**.

Step 2: Inverse Document Frequency (IDF)

Now imagine we have **1,000 books** in the dataset.

- The word “**love**” appears in 800 books → Very common
- The word “**quantum**” appears in only 10 books → Rare

IDF will:

- Give **low weight** to “love”
- Give **high weight** to “quantum”

Final TF-IDF Effect :

Word	TF	IDF	TF-IDF Result
love	High	Low	Medium
quantum	Medium	High	High

Even though “love” appears many times,
it gets **lower importance** because it appears in many books.

But “quantum” gets a **higher score** because:

- It is rare across books
- It helps distinguish science books from romance books

In short :

For example, the word “love” may appear frequently in many books, so its IDF is low and its TF-IDF score is reduced. However, a rare word like “quantum” appears in only a few books, giving it a high IDF and higher TF-IDF score. This makes TF-IDF effective for identifying distinguishing words in documents.

PART 5:

How is cosine similarity calculated? Why is it appropriate for comparing documents represented as TF-IDF vectors?

How is Cosine Similarity Calculated?

Cosine similarity measures the **angle between two vectors**.

For two TF-IDF vectors **A** and **B**:

$$\text{cosine similarity} = \frac{A \cdot B}{\|A\| \|B\|}$$

Where:

- $A \cdot B$ = dot product of the two vectors
- $\|A\|$ = magnitude (length) of vector A
- $\|B\|$ = magnitude of vector B

Interpretation:

Value	Meaning
1	Documents are very similar
0	Documents are very similar
-1	Documents are unrelated

Since TF-IDF values are non-negative, cosine similarity is usually between **0 and 1**.

Why Cosine Similarity is Appropriate for TF-IDF?

1. Why Cosine Similarity is Appropriate for TF-IDF:
 - Longer documents naturally have larger vector magnitudes.
 - Cosine similarity **normalizes by vector length**, so length doesn't bias similarity.
2. Independent of document length:

TF-IDF vectors are:

- High dimensional
- Mostly zeros

Cosine similarity handles sparse vectors efficiently.

3. Works well with sparse vectors:

- It measures **direction**, not magnitude.
- So it compares topic similarity rather than document length.

In short:

Cosine similarity is calculated as the dot product of two vectors divided by the product of their magnitudes. It measures the angle between TF-IDF vectors. It is appropriate because it normalizes document length, works well with sparse high-dimensional vectors, and captures similarity based on word distribution rather than document size.

PART 6:

Discuss scalability challenges when calculating pairwise cosine similarity for many documents. How could Spark help address these challenges?

Scalability Challenges in Pairwise Cosine Similarity:

If there are **N documents**, computing all pairwise similarities requires: $O(N^2)$ comparisons.

Example:

- 1,000 books → 1,000,000 comparisons
- 10,000 books → 100,000,000 comparisons
- 100,000 books → 10 billion comparisons

This becomes:

- Computationally expensive
- Memory intensive
- Slow for large datasets

How Spark Helps:

- Distributed Processing - Spark distributes similarity computation across multiple nodes.
- Parallel Vector Operations - TF-IDF vectors are processed in parallel using Spark MLlib.
- Efficient Data Partitioning - Spark partitions data across executors, enabling parallel comparisons.
- Approximate Methods –

Spark provides:

- **LSH (Locality Sensitive Hashing)**
- **Approximate nearest neighbors**

These reduce computation from $O(N^2)$ to much faster approximate methods.

In short:

Cosine similarity is calculated as the dot product of two TF-IDF vectors divided by the product of their magnitudes. It measures the angle between vectors and ranges from 0 to 1 for TF-IDF representations. It is appropriate for document comparison because it normalizes document length, handles sparse high-dimensional vectors effectively, and captures similarity based on word distribution rather than size. However, computing pairwise cosine similarity for N documents requires $O(N^2)$ comparisons, which becomes computationally expensive for large datasets. Apache Spark addresses this challenge through distributed processing, parallel computation, and approximate similarity techniques such as Locality Sensitive Hashing.

Q-12:

Author Influence Network:

Task:

Construct a simple author influence network based on whether authors' books were released within a certain time window of each other. This is a simplified representation, as true influence is much more complex, but it demonstrates graph-like analysis with Spark.

PART 1- Preprocessing:

STEP 1: Extract Author and Release Year

Extract the author and release_date from the text of each book, similar to what you did in the first question of the previous set. You might need to refine your regular expressions or use more advanced techniques if the format varies significantly.

file_name	title	release_date	language	encoding
10.txt	The King James Bible	March 2, 2011 [EBook #10]	English	ASCII
101.txt	Hacker Crackdown	January, 1994	English	ASCII
102.txt	The Tragedy of Pudd'nhead Wilson	January, 1994	English	ASCII
103.txt	Around the World in 80 Days	May 15, 2008 [EBook #103]	English	ASCII
104.txt	Franklin Delano Roosevelt's First Inaugural Address	May 14, 2008 [EBook #104]	English	ASCII
105.txt	Persuasion	June 5, 2008 [EBook #105]	English	ASCII
106.txt	Jungle Tales of Tarzan	June 5, 2008 [EBook #106]	English	ASCII
107.txt	Far from the Madding Crowd	February, 1994 [eBook #107]	English	ISO-8859-1 (US-ASCII)
108.txt	The Return of Sherlock Holmes	July 8, 2007 [EBook #108]	English	ASCII
109.txt	Renascence and Other Poems	June 19, 2008 [EBook #109]	English	ASCII

only showing top 10 rows

```
>>> from pyspark.sql.functions import regexp_extract, col
>>> author_df = books_df \
    .withColumn(
        "release_date",
        ...     .withColumn(
            ...     ...     "author",
            ...     ...     regexp_extract(..., r"(?i)Author:\s*(.+)", 1)
            ...     ) \ ...
        ...     .withColumn(
            ...     "release_date",
            ...     ...     regexp_extract(col("text"), r"(?i)Release Date:\s*(.+)", 1)
            ...     )
```

Apache Hadoop & MapReduce

```
>>> author_df.show()
26/02/11 18:35:58 WARN PythonRunner: Detected deadlock while completing task 0.0 in stage 3 (TID 428): Attempting to kill Python Worker
+-----+-----+-----+-----+
|file_name|      text|     author| release_date|
+-----+-----+-----+-----+
| 10.txt|The Project Guten...|           |March 2, 2011 [EB...|
|101.txt|The Project Guten...| Bruce Sterling| January, 1994|
|102.txt|The Project Guten...|Mark Twain (Samue...| January, 1994|
|103.txt|The Project Guten...|       Jules Verne|May 15, 2008 [EBo...|
|104.txt|The Project Guten...|Franklin Delano R...|May 14, 2008 [EBo...|
|105.txt|The Project Guten...|   Jane Austen|June 5, 2008 [EBo...|
|106.txt|Project Gutenberg...|Edgar Rice Burroughs|June 5, 2008 [EBo...|
|107.txt|The Project Guten...|    Thomas Hardy|February, 1994 [...|
|108.txt|Project Gutenberg...| Arthur Conan Doyle|July 8, 2007 [EBo...|
|109.txt|Project Gutenberg...|Edna St. Vincent ...|June 19, 2008 [EBo...|
|11.txt|Project Gutenberg...|       Lewis Carroll| March, 1994|
|110.txt|The Project Guten...|    Thomas Hardy|February, 1994 [...|
|111.txt|The Project Guten...|Gene Stratton-Porter|March 8, 2006 [EBo...|
|112.txt|The Project Guten...| Richard McGowan| March, 1994|
|113.txt|The Project Guten...|Frances Hodgson B...|May 15, 2008 [EBo...|
|114.txt|The Project Guten...| Sir John Tenniel|May 27, 2008 [EBo...|
|115.txt|The Project Guten...|U.S. Census of Po...|June 5, 2008 [EBo...|
|117.txt|The Project Guten...|Ludwig van Beethoven|March 16, 2012 [E...|
|118.txt|The Project Guten...|Electronic Fronti...| March, 1994|
| 12.txt|The Project Guten...|Charles Dodgson, ...| February, 1991|
+-----+
only showing top 20 rows
```

STEP 2: Extract Release Year

```
>>> author_df = author_df.withColumn(
...     "release_year",
...     regexp_extract(col("release_date"), r"(\d{4})", 1)
... )
>>> author_df.show
```

```
>>> author_df.show()
26/02/11 18:38:34 WARN PythonRunner: Detected deadlock while completing task 0.0 in stage 4 (TID 428): Attempting to kill Python Worker
+-----+-----+-----+-----+-----+
|file_name|      text|     author| release_date|release_year|
+-----+-----+-----+-----+-----+
| 10.txt|The Project Guten...|           |March 2, 2011 [EB...| 2011|
|101.txt|The Project Guten...| Bruce Sterling| January, 1994| 1994|
|102.txt|The Project Guten...|Mark Twain (Samue...| January, 1994| 1994|
|103.txt|The Project Guten...|       Jules Verne|May 15, 2008 [EBo...| 2008|
|104.txt|The Project Guten...|Franklin Delano R...|May 14, 2008 [EBo...| 2008|
|105.txt|The Project Guten...|   Jane Austen|June 5, 2008 [EBo...| 2008|
|106.txt|Project Gutenberg...|Edgar Rice Burroughs|June 5, 2008 [EBo...| 2008|
|107.txt|The Project Guten...|    Thomas Hardy|February, 1994 [...| 1994|
|108.txt|Project Gutenberg...| Arthur Conan Doyle|July 8, 2007 [EBo...| 2007|
|109.txt|Project Gutenberg...|Edna St. Vincent ...|June 19, 2008 [EBo...| 2008|
|11.txt|Project Gutenberg...|       Lewis Carroll| March, 1994| 1994|
|110.txt|The Project Guten...|    Thomas Hardy|February, 1994 [...| 1994|
|111.txt|The Project Guten...|Gene Stratton-Porter|March 8, 2006 [EBo...| 2006|
|112.txt|The Project Guten...| Richard McGowan| March, 1994| 1994|
|113.txt|The Project Guten...|Frances Hodgson B...|May 15, 2008 [EBo...| 2008|
|114.txt|The Project Guten...| Sir John Tenniel|May 27, 2008 [EBo...| 2008|
|115.txt|The Project Guten...|U.S. Census of Po...|June 5, 2008 [EBo...| 2008|
|117.txt|The Project Guten...|Ludwig van Beethoven|March 16, 2012 [E...| 2012|
|118.txt|The Project Guten...|Electronic Fronti...| March, 1994| 1994|
| 12.txt|The Project Guten...|Charles Dodgson, ...| February, 1991| 1991|
+-----+
only showing top 20 rows
```

```
>>> author_year_df = author_df.select("author", "release_year")
>>> author_year_df = author_year_df.dropDuplicates()
>>> author_year_df = author_year_df.orderBy("release_year")
>>> from pyspark.sql.functions import col
") \
    .filter(
        (col("author").isNotNull()) &
        (col("release_year").isNotNull())
>>> author_year_df = author_df.select("author", "release_year") \
...     .filter(
...         (col("author").isNotNull()) &
...         (col("release_year").isNotNull())
...
>>>
```

```
>>> author_year_df.show()
26/02/11 18:44:44 WARN PythonRunner: Detected error
+-----+-----+
|      author|release_year|
+-----+-----+
|    Bruce Sterling|      1994|
|Mark Twain (Samue...|      1994|
|      Jules Verne|      2008|
|Franklin Delano R...|      2008|
|      Jane Austen|      2008|
|Edgar Rice Burroughs|      2008|
|      Thomas Hardy|      1994|
|Arthur Conan Doyle|      2007|
|Edna St. Vincent ...|      2008|
|      Lewis Carroll|      1994|
|      Thomas Hardy|      1994|
|Gene Stratton-Porter|      2006|
|      Richard McGowan|      1994|
|Frances Hodgson B...|      2008|
|      Sir John Tenniel|      2008|
|U.S. Census of Po...|      2008|
|Ludwig van Beethoven|      2012|
|Electronic Fronti...|      1994|
|Charles Dodgson, ...|      1991|
|      Jane Austen|      1994|
+-----+-----+
only showing top 20 rows
```

PART 2: Construct Influence Network

1. Define an "influence" relationship between two authors if one author released a book within X years of another author's release (e.g., X = 5 years). You can adjust X to explore different influence ranges.

2. Create an RDD or DataFrame representing the edges of this influence network. Each edge should be a tuple or row of the form (author1, author2), indicating that author1 potentially influenced author2 .

Author A influences Author B if: $0 < (\text{yearB} - \text{yearA}) \leq X$

```
>>> from pyspark.sql.functions import abs
>>> a = author_df.alias("a")
>>> b = author_df.alias("b")

>>> X = 5

>>> edges_df = a.join(b, col("a.author") != col("b.author")) \
...     .withColumn("year", year.cast("int")) \
...     .withColumn("a.release_year", a.release_year.cast("int")) \
...     .withColumn("b.release_year", b.release_year.cast("int")) \
...     .filter((col("b.release_year").cast("int") > col("a.release_year").cast("int")) &
...             ((col("b.release_year").cast("int") - col("a.release_year").cast("int")) <= X))
lect(
    co... ) \
...     .select(
...         col("a.author").alias("author1"),
...         col("b.author").alias("author2")
...     )

>>> edges_df.show()
26/02/11 18:52:29 WARN PythonRunner: Det...
er
+-----+-----+
|   author1|      author2|
+-----+-----+
|Bruce Sterling| Project Gutenberg|
|Bruce Sterling| Robert Service|
|Bruce Sterling| Henry James|
|Bruce Sterling| Unknown|
|Bruce Sterling| Andrew Barton 'Ba...|
|Bruce Sterling| Lao-Tze|
|Bruce Sterling| Joseph Conrad|
|Bruce Sterling| Joseph Conrad|
|Bruce Sterling| G. K. Chesterton|
|Bruce Sterling| Theodore Dreiser|
|Bruce Sterling| Giuseppe Salza|
|Bruce Sterling| United States|
|Bruce Sterling| Robert Harris|
|Bruce Sterling| Arthur Conan Doyle|
|Bruce Sterling| Omar Khayyam|
|Bruce Sterling| Jean Clottes|
|Bruce Sterling| Michael Hart|
|Bruce Sterling| United States|
|Bruce Sterling| Dan Scavelle|
|Bruce Sterling| United States|
+-----+
only showing top 20 rows
```

PART 3: Compute In-Degree and Out-Degree

1. Calculate the "in-degree" and "out-degree" of each author in the network. In-degree represents the number of authors who potentially influenced them, and out-degree represents the number of authors they potentially influenced.

Out-Degree (how many influenced):

```
out_degree = edges_df.groupBy("author1") \
    .count() \
    .withColumnRenamed("count", "out_degree")
```

```
>>> out_degree = edges_df.groupBy("author1") \
...     .count() \
...     .withColumnRenamed("count", "out_degree")
>>> out_degree.show()
+-----+-----+
|      author1|out_degree|
+-----+-----+
|      Henry Lawson|      13|
|      Lewis Carroll|     255|
|      Robert Nemiroff|      26|
|The National Atom...|      13|
|      Henry H. Snelling|      13|
|      Robert W. Service|     119|
|      Jules Verne|     255|
|      Mark E. Laxer|     113|
|David Herbert Law...|     177|
|      James Branch Cabell|      13|
|United States Bur...|     286|
|      Leo Tolstoy|     177|
|      Samuel D. Humphrey|      13|
|Edna St. Vincent ...|      13|
|      Adam Lindsay Gordon|      13|
|Franklin Delano R...|      13|
|          Plato|      13|
|Charles Dodgson, ...|     162|
|      H. G. Wells|     432|
|Edgar Rice Burroughs|     720|
+-----+-----+
only showing top 20 rows
```

In-Degree (how many influenced them):

```
in_degree = edges_df.groupBy("author2") \
    .count() \
    .withColumnRenamed("count", "in_degree")
```

```
>>> in_degree = edges_df.groupBy("author2") \
...     .count() \
...     .withColumnRenamed("count", "in_degree")
>>> in_degree.show()
+-----+-----+
|      author2|in_degree|
+-----+-----+
|      Henry Lawson|      66|
|      Lewis Carroll|      45|
|      Franz Josef Haydn|      56|
|      Robert Nemiroff|     132|
|The National Atom...|      66|
|      Henry H. Snelling|      66|
|      Robert W. Service|     178|
|      Jules Verne|     130|
|      James Branch Cabell|      66|
|      Mark E. Laxer|      39|
|David Herbert Law...|      12|
|      Giuseppe Salza|      56|
|United States Bur...|      14|
|      Joyce Kilmer|      56|
|      Leo Tolstoy|      12|
|      Samuel D. Humphrey|      66|
|      Geoffrey Chaucer|      56|
|Edna St. Vincent ...|      66|
|      Adam Lindsay Gordon|      66|
|Franklin Delano R...|      66|
+-----+
only showing top 20 rows
```

2. Identify the top 5 authors with the highest in-degree and the top 5 authors with the highest out-degree.

Top 5 by Out-Degree:

```
>>> out_degree.orderBy(col("out_degree").desc()).show(5)
+-----+-----+
|      author1|out_degree|
+-----+-----+
| Thomas Hardy|      828|
|Robert Louis Stev...|  817|
|Edgar Rice Burroughs| 720|
|      Henry James|  613|
|Frances Hodgson B...| 551|
+-----+-----+
only showing top 5 rows
```

Top 5 by In-Degree:

```
>>> in_degree.orderBy(col("in_degree").desc()).show(5)
+-----+-----+
|      author2|in_degree|
+-----+-----+
|Robert Louis Stev...| 1911|
|Edgar Rice Burroughs| 941|
|      Arthur Conan Doyle| 615|
|      L. Frank Baum| 552|
|Richard Harding D...| 512|
+-----+-----+
only showing top 5 rows
```

PART 4:

Discuss how you chose to represent the influence network in Spark (e.g., RDD of tuples, DataFrame). What are the advantages and disadvantages of your chosen representation?

I represented the author influence network as a **Spark DataFrame** containing directed edges of the form: (author1, author2).

Where:

- author1 → potentially influenced
- author2 → potentially influenced by
- Each row represents one directed influence relationship.

Why I Chose DataFrame (Instead of RDD)?

Advantages:

Apache Hadoop & MapReduce

1. Optimized Execution- Spark SQL optimizer (Catalyst) improves join and aggregation performance.
2. Easy Degree Computation - In-degree and out-degree can be computed easily using: `groupBy().count()`
3. Scalable –
 - Works efficiently for large datasets.
 - Supports distributed joins.
4. Readable & Structured –
 - Schema-based.
 - Easier debugging and filtering.

Disadvantages:

1. Not a Native Graph Structure - No built-in graph traversal (e.g., PageRank, shortest path) unless using GraphFrames.
2. Self-Join Can Be Expensive - Influence detection requires comparing authors, which may become costly for large datasets.
3. Limited Graph Semantics - Unlike GraphX or GraphFrames, relationships are not automatically treated as graph objects.

Why Not RDD?

RDD of tuples like:

`("Jane Austen", "Charles Dickens")`

would work, but:

- No query optimization
- Harder aggregation
- More manual coding
- Slower for large joins

I represented the influence network as a Spark DataFrame of directed edges (`author1, author2`). This approach enables efficient distributed joins and easy degree calculations using group-by operations. While it is scalable and optimized, it lacks native graph-processing capabilities compared to GraphFrames or GraphX.

PART 5:

How does the choice of the time window (X) affect the structure of the influence network?
 What are the limitations of this simplified definition of influence?

1. How does the choice of time window (X years) affect the network?

Assume we define influence as:

Author A influences Author B if B published within **X years** after A.

If X is Small (e.g., 2–5 years)

- Fewer influence edges

Apache Hadoop & MapReduce

- Network becomes **sparser**
- Only very close temporal relationships are captured
- Likely lower in-degree and out-degree
- May miss long-term literary influence

Example:

If Jane Austen (1813) and Dickens (1837) → 24 years apart

With $X = 5 \rightarrow$ no edge

With $X = 30 \rightarrow$ edge appears

If X is Large (e.g., 30–50 years)

- Many more edges
- Network becomes **denser**
- Authors appear highly connected
- May create unrealistic influence links
- Can inflate in-degree/out-degree

Too large $X \rightarrow$ almost everyone influences everyone.

So the time window directly controls:

- Edge density
- Connectivity
- Degree distribution
- Centrality of authors

2. Limitations of This Simplified Influence Definition

This definition is **very naive**. It assumes:

Time proximity = influence

But in reality, influence is much more complex.

Major Limitations:

1. No Content Similarity - Two authors close in time may write completely different genres.
2. No Citation or Historical Evidence –

It doesn't check whether:

- Author B read Author A
 - Author B referenced Author A
 - There is stylistic similarity
3. Ignores Long-Term Influence - Some authors influence writers decades later (e.g., Shakespeare).
 4. False Positives - Just because books were published close in time doesn't mean influence exists.

5. Ignores Multiple Works - An author may publish many books over time — using only one release date oversimplifies.

Better Real-World Approach:

A stronger influence model could combine:

- Time proximity
- Text similarity (TF-IDF cosine similarity)
- Citation/reference analysis
- Genre similarity
- Historical records

In short:

A smaller time window produces a sparse network with fewer influence edges, while a larger window creates a denser network with many connections. However, defining influence purely based on temporal proximity is limited because it ignores textual similarity, citations, genre, and historical evidence. Time alone does not guarantee actual influence.